

MedicScribe

A private medical scribe for Malaysian clinics

by Sunkarilabs

Data Sovereignty · Zero-Cloud Liability · Specialized Medical Intelligence

What it is: a tool that listens to your consultation and writes a structured clinical note for you to review, edit, and export — so you spend the visit with the patient, not the keyboard.

Status: working and field-tested today. The screenshots in this document are from the live app.


The promise: your patients' data is **sovereign** — it stays on hardware you control and never touches the cloud. No third party ever holds, sees, or could leak it.

How it is delivered: a sealed, headless "black box" AI server that runs on its own — locked down, with nothing to manage and nothing to tinker with.

1. What it does today

You press one button at the start of the consultation and one button at the end. About a few seconds after you stop, a clean SOAP note appears on your screen. You read it, fix anything you want, and share or export it.

The app is built on speech-recognition models tuned for Malaysian multilingual consultations — it handles English, Malay, Mandarin and Tamil, plus local **dialects and everyday slang (Manglish)**, including switching between them in the same sentence — and still produces an **English** note.



1. **Start** — one button, and a banner reminding you the data is private, not the cloud.

2. **Recording** — the button becomes Stop. Nothing else to manage.

3. **The note** — a ready SOAP note you can edit and Share.

4. **Privacy trail** — proof of exactly what happened to the data.

In the real example above, the system captured the patient's complaint and a detailed medication plan from a two-minute conversation, and the note was ready a few seconds after Stop. The privacy trail shows the audio was **deleted on the server**, the note is **export-only** — **never stored**, and everything is **wiped when you tap Done**.

2. How it works (in plain terms)

The phone or tablet does **not** do the thinking. It only records your voice and shows you the result. The actual work is done by two AI models running on a separate computer — the **server**:

1. **You record** — the app captures the consultation audio.
2. **Speech → text** — the first model turns the spoken consultation into written text, translating any mix of languages into English along the way.
3. **Text → note** — the second model reads that text and organises it into a structured SOAP note.
4. **You review** — the note returns to your screen to edit, share, and export.

Two buttons for you. Two models doing the work behind the scenes.

3. Data Sovereignty — why it never touches the cloud

Most medical scribes on the market send the recording to a company's servers on the internet. For patient data, that is a permanent liability. MedicScribe runs on **one dedicated machine you control** — either a private Sunkarilabs server (evaluation) or a sealed server installed in your own clinic.

- **Data sovereignty:** the audio and the note live only on your hardware. No cloud account, no foreign data centre, no vendor holding your records.
- **Zero-cloud liability:** nothing leaves the building, so there is no cloud breach surface to be liable for. This can be guaranteed contractually (see the PDPA Audit Trail add-on).
- **Offline by design:** a clinic server keeps working even if your internet is down.
- **No lock-in:** built on open-source models — no per-record cloud fees, no subscription you can't leave.

4. Why the hardware matters (and why a bigger one costs more)

The AI models run on a **graphics card (GPU)** inside the server. This single component decides how good and how fast the notes are:

- A GPU has a **fixed amount of working memory** — think of it as a fixed-size desk. Each model must fit entirely on that desk to run **at all**. If a model is bigger than the desk, it will not start.
- A **smarter, more accurate model is also a bigger model**, so it needs a **bigger desk** (more GPU memory).
- **More doctors recording at once** means more jobs sharing the desk — again, more room needed.

This is why the on-site standard uses a **24 GB GPU** for 1–2 doctors, and the multi-doctor tier uses a multi-GPU architecture for 3–5 doctors running in parallel.

Higher accuracy or more doctors → a bigger GPU → higher hardware cost. Hardware is the lower-cost lever; bespoke model work is the premium one. We size the hardware to the clinic.

5. A sealed "black box" — reliability, SLA & uptime

The server is delivered **headless** — no screen, no desktop, nothing on it to open or change. Everything runs automatically inside the box. This is deliberate:

- **All the graphics-card memory goes to the models.** A screen or desktop on the server would quietly use part of that same memory; on a smaller server that can be enough to tip it over and stop a note from running. Headless keeps every bit of the card free for making notes.
- **Full, exclusive control of the card.** Nothing else runs, so no stray program competes for memory and interrupts a note.
- **No tinkering, no leaks.** Nothing for anyone to fiddle with, mis-configure, or copy data out of. A sealed system is a safe system.
- **Engineered for uptime.** Because the system is locked down with no software clutter, it needs virtually zero manual maintenance — built for **99.9% uptime**.

If it ever does fail, you lose nothing. Notes are export-only and nothing is half-saved, so there is no data to lose — you simply resume the moment the server is back.

Support & SLA: your monthly licence is a direct line to the engineer who built the system — not a call-centre queue. Issues are attended to personally, on weekends included. For on-site clinics, an optional **Cold-Swap hardware SLA** guarantees a physical replacement server within 24 hours (see add-ons).

6. Packages & pricing

Three tiers, from a low-commitment pilot to a multi-doctor clinic deployment. All figures are in Malaysian Ringgit and confirmed against the final specification at quotation.

Tier	What is included	Rate (MYR)
1. Private Cloud Evaluation	A secure, encrypted private link to a dedicated Sunkarilabs server. Strictly a single-doctor pilot, 3 months maximum — to prove the value before committing to on-site hardware.	RM 1,500 / month + RM 3,000 one-time setup
2. Local On-Site "Black Box" RECOMMENDED	A sealed, headless AI server physically installed inside your clinic . 100% offline — zero internet or cloud reliance. Supports 1–2 doctors . The enterprise standard.	RM 15,000 hardware (one-time) + RM 5,000 installation + RM 650 / month licence & SLA
3. Multi-Doctor Clinic	A high-spec multi-GPU server built to run heavy parallel inference for 3–5 doctors simultaneously , fully on-site and offline.	RM 28,000 hardware (one-time) + RM 8,000 installation + RM 1,200 / month licence & SLA

The monthly licence & SLA is your retainer: priority access to the engineer who built the system, software updates, and the reliability guarantee — not a charge for "updates" alone.

7. Premium add-ons

These extend the platform with specialised execution that generic, mass-market AI tools cannot provide.

A · Specialty Vocabulary Tuning — Specialized Medical Intelligence

Off-the-shelf AI models fail at specific clinical jargon. We custom-tune the underlying language models for your medical branch's focus. *Example — Aesthetic & Dermatology Pack*: custom weights trained to recognise laser types (Picosecond, Q-Switched Nd:YAG), conditions (Melasma, Post-Inflammatory Hyperpigmentation), and premium skincare brand names.

From RM 12,000 one-time, per specialised model.

B · Premium Hardware & Hardware SLA

Black-Box + Mic Bundle: a medical-grade, highly sensitive directional wireless microphone paired with a custom-anodised, silent mini-ITX server chassis — sleek on a clinic counter, and it eliminates tablet-microphone distortion.

RM 3,500 per consultation room.

Cold-Swap Guarantee (Hardware SLA): if an on-site server hits a hardware fault, Sunkarilabs physically delivers and swaps a temporary replacement server within **24 hours, weekends included**.

+ RM 250 / month on the base maintenance tier.

C · Compliance & Custom Integration

PDPA Enterprise Audit Trail: a signed compliance certificate and a localised cryptographic logging system that proves no patient audio or text was retained on-site or leaked.

Included in Tier 2 & 3 · RM 2,000 one-time for Tier 1.

Custom EMR / CMS Webhook Integration: a secure local API endpoint that pushes generated SOAP notes straight into your clinic's existing medical-records software with one click.

From RM 5,000, quoted on your software's complexity.

8. iPad and Android — which device

The app can be supplied on **both an iPad and an Android tablet**. They differ only in how soon each is ready:

- **Android:** ready now — this is the field-tested app in the screenshots.
- **iPad / iPhone (iOS):** needs additional development time, because Apple's App Store rules are strict. The practical route is to run MedicScribe as a **secure local web app (HTTPS)** served by your own clinic server and opened in the iPad's browser — no App Store needed.

For the iPad/iPhone route to be trusted and encrypted, each device needs a small **security certificate ("CA") installed once**. After that one-time step it behaves like a normal app, and stays **fully local — still no cloud**. The iPad is only the *recorder*; the AI work happens on the server either way. App layout changes and the iOS build are delivered as a quoted app-design upgrade.

9. Privacy, compliance & PDPA

- **No cloud.** Audio and notes stay on hardware you control — full data sovereignty.
- **Sealed server.** Headless and locked down — nothing to tinker with, nothing to leak.
- **Zero data retention** — the server only runs the models and transcribes; nothing is kept. Guaranteed by contract, provable with the PDPA Enterprise Audit Trail.
- **Audio deleted** on the server right after the note is generated (the app shows the confirmation).
- **Notes are export-only** — never stored on the server.
- **Everything wiped** when you tap Done, with an on-screen audit trail each time.
- **No data loss** if the server ever fails — nothing is half-saved; just resume.
- **Built for Malaysia's PDPA 2010** — open-source models, no lock-in, no per-record cloud fees.

Get in touch

Surenther Saravanan — Sunkarilabs
Phone / WhatsApp: **014-635 7916**
Email: surenthersaravanan@gmail.com

MedicScribe by Sunkarilabs — Data Sovereignty · Zero-Cloud Liability · Specialized Medical Intelligence.